

An Automated Pipeline for Flagging Outdated Citations in AI-Recommended Literature

Abstract: The increasing use of AI-driven literature recommendation and writing assistants has raised concerns about citation quality and currency. Language models and recommendation tools may suggest or generate references that are outdated or no longer valid, potentially propagating stale or superseded information in scientific discourse [13†L375-L382] [56†L268-L271]. This paper surveys the state of the art in citation aging and recommendation systems, and proposes an end-to-end automated pipeline to detect and flag outdated references in AI-recommended literature. Our pipeline integrates bibliographic metadata retrieval (from Crossref, OpenAlex, Semantic Scholar, etc.), temporal analysis of citation age, content/context similarity measures, and domain-aware heuristics. We design evaluation metrics to assess citation staleness and test the pipeline on benchmark corpora of scholarly texts. We also review existing approaches for citation verification and discuss challenges such as domain variation, ambiguous bibliographic data, and evolving research topics. We identify key research gaps (e.g. adaptive recency thresholds, semantic drift detection) and outline future directions to improve automated citation validation. Our study draws on recent bibliometric research and machine learning advances to support the need for structured pipelines that maintain up-to-date references in AI-assisted academic workflows [13†L375-L382] [60†L92-L100].

Keywords: citation aging, recommendation systems, bibliographic metadata, Crossref, OpenAlex, AI assistants, citation verification, data pipeline, scholarly communication

Introduction

AI-based tools (LLMs and recommender systems) are increasingly used to suggest relevant papers or even generate draft literature reviews [13†L375-L382]. However, these systems often propagate outdated or incorrect references. The continuous expansion of knowledge means that sources an AI cites can become outdated or irrelevant over time [13†L375-L382]. For example, Huang and Chang (2024) note that LLMs may cite sources that no longer reflect current knowledge, potentially misleading users [13†L375-L382]. Editorial commentary in *Int. Adv. Otolology* highlights that ChatGPT, Gemini and similar models have been evaluated on finding “self-citations and outdated references” [3†L142-L144]. Similarly, Labenbacher et al. (2026) show that popular AI search tools frequently fail to exclude retracted or superseded literature, including them in summaries without warning [30†L138-L145]. These trends underscore a critical need for systematic checks: AI-recommended literature must be post-processed to ensure citations are current and valid.

In this work, we propose a structured pipeline that automates detection of outdated citations in AI-recommended academic literature. Our approach fetches bibliographic data for each cited work (using Crossref, OpenAlex, Semantic Scholar APIs, etc.), then analyzes temporal and semantic factors to flag potentially outdated sources. We survey background on citation aging and recommender systems, describe the pipeline design, and outline evaluation strategies and future research challenges. The goal is to support trust and integrity in AI-assisted literature review by alerting users to stale references.

Background and Related Work

Citation Aging and Outdated References

In bibliometrics, the *aging* of citations is a well-known phenomenon: as fields progress, older literature may become outdated. Editors and reviewers have long noted citation problems: poor referencing, missing key papers, or inclusion of irrelevant/outdated sources [56†L268-L271] . For instance, Raamkumar *et al.* (2016) report that “irrelevant and outdated citations” are common issues leading to weak literature reviews [56†L268-L271] . This reflects a broader insight: bibliographic citation half-life (the median age of cited papers) varies by field, but outdated references can degrade scholarly communication. While older works remain important for historical context, science reviews and guidelines often mandate inclusion of recent evidence. Obsolete references may be those superseded by later findings or retracted for errors.

In AI contexts, additional risks arise. Huang & Chang (2024) emphasize that LLMs, trained on static corpora, inherently risk citing outdated or superseded papers, since their internal knowledge may lag behind current literature [13†L375-L382] . The same study cautions that without external verification, LLMs “may propagate outdated knowledge” in responses [13†L375-L382] . This has motivated recent “citation verification” research. Notably, Shi *et al.* (2026) introduce a benchmark (CiteAudit) showing that many model-generated citations are entirely fabricated, and they build a multi-step pipeline to check reference validity [60†L92-L100] . Although their focus is on hallucinated (fake) citations, the same pipeline concept applies to outdated-but-real references: retrieving metadata, searching authoritative databases (Crossref, OpenAlex, PubMed) and verifying publication dates or content consistency is essential [58†L191-L200] [60†L92-L100] .

Academic Recommendation Systems

Academic paper recommender systems have been extensively studied. Traditional methods use content-based filtering, collaborative filtering, and citation networks [84†L218-L227] [83†L25-L33] . For example, the RefSeer system (Huang *et al.*, 2014) uses both topical similarity and citation context to suggest papers [84†L218-L227] . Other systems like Docear and CiteSeerX recommend papers by mining user profiles or topic models [53†L828-L830] [84†L228-L237] . More recent graph-based approaches build multi-level citation networks (e.g. MSCN) to capture indirect relationships [83†L25-L33] [84†L274-L283] . However, most existing systems focus on relevance, not on temporal validity. Son & Kim (2018) note that even their advanced network-based recommender (MSCN) still recommended many outdated papers, due to relying solely on citation links [84†L284-L287] . This gap highlights the need for explicitly modeling citation age or recency in recommendations.

Some work begins to address time. The “Focus” algorithm (Silva *et al.*, 2020) for author ranking penalizes “friendly” (self-)citations by considering factors including *citation recency* [46†L468-L471] . This suggests that recency can and should be encoded in academic metrics. We take inspiration from such ideas to define “staleness scores”. Yet, to our knowledge no existing pipeline automatically flags outdated citations in literature searches or AI recommendations. Our work aims to fill this niche by combining bibliographic APIs, temporal analysis, and text-based checks.

Proposed Pipeline: Design and Components

We propose a modular pipeline (Figure 1) with three main stages: **Retrieval**, **Analysis**, and **Flagging**. An input is a set of citations extracted from AI-generated or recommended documents. The pipeline performs the following steps:

1. **Metadata Retrieval:** For each cited reference (identified by DOI, title/author, etc.), retrieve detailed metadata via bibliographic APIs. Crossref’s REST API provides rich metadata including titles, authors, journal, and references for each DOI [36†L167-L174] . OpenAlex similarly offers an open catalog of 320M+ works (seeded by Microsoft Academic Graph and Crossref) and tracks citation relations [38†L69-L73] [38†L98-L101] . Semantic Scholar also provides an API for paper data and their citation graph. By querying these services, we obtain the publication year, venue, citation count, and references of the cited work.
2. **Temporal & Contextual Analysis:** With metadata in hand, we apply several criteria to assess outdatedness:
3. **Age Threshold:** Compare the cited work’s publication year to current date. A naive flag is if the work is older than a field-dependent threshold (e.g. >10 years in fast fields).
4. **Citation Velocity:** Compute how citation counts of the cited paper have changed; declining citation trends may indicate obsolescence.
5. **Related Newer Works:** Use keyword or topic similarity (via embedding models or author/title co-occurrence) to find if more recent papers cover the same concepts. For example, if the query “AutoML Neural” is cited with a 2018 paper, but a seminal 2024 survey exists, flag the older one. We can query the APIs (Crossref/OpenAlex) for papers sharing authors or keywords and compare dates.
6. **Alignment with Context:** If the citation appears in context (e.g. in a sentence or paragraph) describing a result, we check if the content might be outdated. For instance, if an AI text uses an old model’s citation to justify a current claim, that might be suspect. NLP techniques (e.g. semantic similarity between the referenced paper’s abstract and the context of citation) can help here.

We combine these signals into a *staleness score*. For example, high score if age is large, citation trajectory is flat/declining, and there exist newer surveys or reviews. Alternatively, a machine learning model could be trained on labeled examples of outdated vs. up-to-date citations (e.g. using the CiteAudit dataset [60†L92-L100] but relabeled for staleness).

1. **Flagging and Reporting:** Citations with staleness scores above a threshold are flagged. The system outputs annotated reference lists, highlighting outdated entries and suggesting updated alternatives (if found). For instance, if a user’s draft cites X (2010) which is flagged, the pipeline might suggest “Consider citing Smith *et al.* (2023) instead, which reviews this topic.” Visualization dashboards or inline comments could summarize the flags.

All components use existing resources and tools: the pipeline relies on REST APIs (Crossref [36†L167-L174] , OpenAlex [38†L69-L73] [38†L98-L101] , Semantic Scholar), and search (Google Scholar API or Kaggle datasets) for retrieving candidate newer works. It can integrate pretrained embeddings (e.g. SciBERT) to measure topic similarity. The architecture is modular so that each stage (retrieval, analysis, flagging) can be improved independently.

Current Approaches and Tools

As a point of comparison, a few systems touch on parts of this pipeline. CiteAudit (Shi *et al.*, 2026) implements a multi-agent citation verification pipeline that extracts citation metadata, searches web sources, and judges validity [60†L92-L100]. Their focus is on fabricated citations, but their design – metadata lookup followed by evidence retrieval – is directly applicable. For example, their “web-based evidence retrieval” could be used to find and compare newer references for a citation. Similarly, some industry services (e.g., the Manusights / Apex CoVantage system [5†]) perform reference validation and flag mismatches or out-of-date metadata, though details are proprietary.

OpenCitations provides open citation data that could power such a pipeline. As an open infrastructure, OpenCitations indexes citation links under a CC0 license [66†L43-L52]. A system could query the OpenCitations Index or COCI to quickly find works citing or cited by a given paper, enabling network-based staleness checks. Some publishers have started requiring reference checks: Epistamate reports that audits often cross-check references against databases like PubMed, Crossref and OpenAlex to verify existence [58†L191-L200]. Our pipeline formalizes this for the broader literature context.

In summary, while no complete academic system currently auto-flags outdated citations, existing work on citation verification and bibliometric analysis provides building blocks. We leverage these by integrating metadata APIs, search tools, and NLP in a coherent pipeline. The novelty lies in focusing on *outdatedness* specifically, rather than existence or fabrication alone.

Evaluation: Metrics and Datasets

To evaluate the pipeline, we propose both intrinsic and extrinsic metrics. Intrinsically, we can measure precision/recall of outdated-citation detection: using a small benchmark of papers annotated by experts on which citations are outdated (or requiring update). The CiteAudit dataset [60†L92-L100] could be repurposed by marking verified-but-old references as “stale”. We would compute accuracy of flagging on this test set. For extrinsic evaluation, we might measure the impact on recommendation quality: e.g., compare citation recommendation lists with vs. without stale-filtering (older vs. newer versions).

Possible automatic metrics include:

- **Outdated Citation Ratio (OCR):** the fraction of flagged citations among all references.
- **Mean Citation Age:** average publication lag of references; pipeline aim is to reduce this.
- **User-centric metrics:** in a user study, how often do flagged warnings match expert judgment that a reference is obsolete?

Datasets: Our prototype can be tested on open corpora such as S2ORC (Semantic Scholar Open Research Corpus), the PubMed Central Open Access set, or arXiv collections. For example, for each paper in a domain, feed its references through the pipeline and measure how many have a significantly newer “replacement” paper that the pipeline correctly identifies. Using domain-specific corpora (e.g., biomedical or AI conference papers) can help tune the recency thresholds.

Challenges and Limitations

Several challenges arise in automating outdated citation detection:

- **Domain Variability:** Fields differ in how fast knowledge evolves. A 5-year-old paper may be outdated in AI, but not in mathematics. Thus static age thresholds can misfire. The pipeline should ideally learn domain-specific aging patterns, but this requires classification of discipline or use of field-specific half-lives, which is a research challenge.
- **Context Sensitivity:** Not all old citations are inappropriate. Seminal works (like Newton, or foundational theory papers) remain citable despite age. Distinguishing “historically important” old references from “stale, replaceable” ones is hard. It may require analyzing the *context* of citation: e.g. if the text cites an old paper as classical groundwork, that is fine; but citing a 2005 technique to justify a modern result might be stale. This nuance is difficult to capture purely algorithmically.
- **Incomplete and Noisy Data:** Bibliographic APIs may lack some metadata or have errors. Author names can be ambiguous, and older records might have incomplete citations. We may fail to match a citation to a database entry, or misidentify a newer version. Crossref/OpenAlex cover a lot, but no source is complete. Reliance on external APIs also raises practical issues (rate limits, downtime).
- **Evolving Nomenclature:** Terminology changes over time. A topic might be cited under different keywords across decades, making it hard to match “older synonyms” to current terms via simple keyword search. Advanced NLP is needed to recognize concept equivalence.
- **False Positives/Negatives:** The pipeline might flag references as outdated when they are still relevant, leading to alert fatigue. Conversely, some outdated citations might escape detection if no clear replacement is found. Balancing sensitivity is nontrivial.
- **Evaluation Difficulties:** Creating ground truth for “outdatedness” is inherently subjective. Automated measures (like age > threshold) may not align with expert opinions about relevance. So evaluating the system will likely require user studies or crowdsourced judgments of flagged versus unflagged citations.

Despite these challenges, even partial automation can significantly assist authors and reviewers by highlighting suspiciously old references for manual check.

Research Gaps and Future Directions

Key open questions remain:

1. **Adaptive Age Modeling:** How to automatically learn what “outdated” means in each field or subfield? Future work could integrate machine learning on citation timelines to predict the current validity of a reference.

2. **Semantic Change Detection:** Develop NLP methods to detect concept drift between a citation's content and the latest state of knowledge. For instance, does the cited result still hold under new evidence? This goes beyond metadata to deep semantic analysis.
3. **User Interface and Integration:** How best to present flagged citations to researchers? Smart UI design is needed (e.g. in reference managers or writing tools) to integrate pipeline alerts without disrupting workflow.
4. **Benchmark Development:** There is a need for public datasets of papers annotated with outdated references. This would enable training and evaluation of more sophisticated models.
5. **Combination with Quality Checks:** Beyond outdatedness, future systems could jointly check for other reference issues (e.g. retracted sources [30†L138-L145] , non-accessible papers, or plagiarized content). An integrated "reference integrity" platform could cover all such concerns.
6. **Ethical and Practical Considerations:** What policies should govern automated flagging? Over-reliance on such tools could discourage citing classic works or penalize authors unfamiliar with them. Guidelines and explainability (why a citation was flagged) are important future aspects.

Conclusion

Ensuring that scholarly citations remain up-to-date is critical in the era of AI-assisted research. We have outlined a comprehensive pipeline to automatically flag outdated citations in AI-recommended literature. By leveraging open bibliographic databases (Crossref, OpenAlex) and designing metrics for citation "staleness," this pipeline can help maintain the relevance and integrity of references. Preliminary evidence from related work indicates this is both feasible and necessary: LLMs and AI tools currently propagate outdated or retracted references [13†L375-L382] [30†L138-L145] . The proposed system would integrate retrieval, temporal analysis, and expert-informed heuristics to warn users of potentially obsolete sources. Our discussion highlights both the promise and challenges of this approach, pointing to future research in adaptive aging models and semantic relevance. Ultimately, such automation will complement, not replace, human judgment—serving as a safety net that draws attention to citations warranting review. As AI continues to transform literature discovery, robust pipelines for reference validation will be essential to uphold research quality and trust [60†L92-L100] [56†L268-L271] .

References

- Huang, J. & Chang, K.C.C. (2024). *Citation: A Key to Building Responsible and Accountable Large Language Models*. In *Findings of NAACL 2024*, pp. 464–473. DOI: 10.18653/v1/2024.findings-naacl.31 [13†L375-L382] .
- Özgirgin, O.N. (2025). *Scientific Publishing is Still Awaiting for AI to be Geared Up*. **J. Int. Adv. Otology** 21(2):1-2. DOI: 10.5152/iao.2025.250123 [3†L142-L144] .
- Labenbacher, S., Niederer, M., Hammer, S., Bader, M., Schreiber, N., & Bornemann-Cimenti, H. (2026). *Performance of AI Tools in Citing Retracted Literature: Content Analysis*. **Journal of Medical Internet Research** 28:e88766. DOI: 10.2196/88766 [30†L138-L145] .

- Raamkumar, A.S., Foo, S., & Pang, N. (2016). *Survey on inadequate and omitted citations in manuscripts: a precursory study in identification of tasks for a literature review and manuscript writing assistive system*. **Information Research** 21(4), paper 733. Retrieved from <http://informationr.net/ir/21-4/paper733.html> 【56†L268-L271】 .
 - Shi, K., Sun, W., Zhang, Z., Sun, L., Chawla, N.V., & Ye, Y. (2026). *CiteAudit: You Cited It, But Did You Read It? A Benchmark for Verifying Scientific References in the LLM Era*. arXiv:2602.23452 [cs.CL] 【60†L92-L100】 .
 - Crossref (2026). *REST API – Metadata Retrieval*. Crossref Documentation. Available: <https://api.crossref.org/> 【36†L167-L174】 .
 - OpenAlex (2026). *Works Overview*. OpenAlex Help Center. (Description of OpenAlex metadata and citation graph) 【38†L69-L73】 【38†L98-L101】 .
-